

Name: Nikhil Arora

Loading Packages

```
In [1]: 1 import pandas as pd
2 import numpy as np
3 import tensorflow as tf
4 import matplotlib.pyplot as plt
5
6 from keras.models import Sequential
7 from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D
8 from keras.utils import to_categorical
9 from keras.optimizers import SGD
```

Using TensorFlow backend.

Importing and Reshaping the Dataset

```
In [2]: 1 from keras.datasets import fashion_mnist
2 (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
3
4 # reshape dataset to have a single channel
5 x_train = x_train.reshape((x_train.shape[0], 28, 28, 1))
6 x_test = x_test.reshape((x_test.shape[0], 28, 28, 1))
7
8 # Normalizing the pixel values to be in the range [0, 1]
9 # by dividing them by 255
10 x_train = x_train/255.0
11 x_test = x_test/255.0
12
13 # Converting the labels into one-hot vectors
14 y_train = to_categorical(y_train, num_classes = 10)
15 y_test = to_categorical(y_test, num_classes = 10)
```

Task 1: Build a neural network without convolutional layers to do the classification task

Step 1: Defining a Model Architecture

```
In [3]: 1 model_T1 = Sequential()
2
3 #Input Layer
4 model_T1.add(Flatten(input_shape=(28, 28, 1)))
5
6 #Hidden Layers
7 model_T1.add(Dense(128, activation='relu'))
8 model_T1.add(Dense(64, activation='relu'))
9
10 #Output Layer
11 model_T1.add(Dense(10, activation='softmax'))
```

WARNING:tensorflow:From C:\Users\nikhi\anaconda3\envs\NeuralNetwork\lib\site-packages\tensorflow_core\python\ops\resource_variable_ops.py:1630: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass *_constraint arguments to layers.

Step 2: Compiling the model

```
In [4]: 1 # Defining Optimizer, Loss function and evaluation metrics
2 model_T1.compile(loss='categorical_crossentropy', optimizer='adam', met
```

```
In [5]: 1 # Model Structure
2 model_T1.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 784)	0
dense_1 (Dense)	(None, 128)	100480
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 10)	650

=====
Total params: 109,386
Trainable params: 109,386
Non-trainable params: 0

Step 3: Training the Model

```
In [6]: 1 # Setting a random state so that comparing the evaluation metrics is po
        2 tf.set_random_seed(1)
        3 np.random.seed(1)
        4
        5 model_T1.fit(x_train, y_train, epochs=10, batch_size=32, validation_dat
```

WARNING:tensorflow:From C:\Users\nikhi\AppData\Local\Temp\ipykernel_9356\3595446061.py:2: The name tf.set_random_seed is deprecated. Please use tf.compat.v1.set_random_seed instead.

WARNING:tensorflow:From C:\Users\nikhi\anaconda3\envs\NeuralNetwork\lib\site-packages\keras\backend\tensorflow_backend.py:422: The name tf.global_variables is deprecated. Please use tf.compat.v1.global_variables instead.

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/10
60000/60000 [=====] - 5s 82us/step - loss: 0.4904
- accuracy: 0.8259 - val_loss: 0.4505 - val_accuracy: 0.8413
Epoch 2/10
60000/60000 [=====] - 5s 88us/step - loss: 0.3691
- accuracy: 0.8654 - val_loss: 0.3945 - val_accuracy: 0.8534
Epoch 3/10
60000/60000 [=====] - 5s 81us/step - loss: 0.3340
- accuracy: 0.8770 - val_loss: 0.3775 - val_accuracy: 0.8679
Epoch 4/10
60000/60000 [=====] - 5s 81us/step - loss: 0.3108
- accuracy: 0.8848 - val_loss: 0.3574 - val_accuracy: 0.8713
Epoch 5/10
60000/60000 [=====] - 5s 80us/step - loss: 0.2935
- accuracy: 0.8900 - val_loss: 0.3512 - val_accuracy: 0.8673
Epoch 6/10
60000/60000 [=====] - 5s 79us/step - loss: 0.2771
- accuracy: 0.8977 - val_loss: 0.3364 - val_accuracy: 0.8815
Epoch 7/10
60000/60000 [=====] - 5s 82us/step - loss: 0.2668
- accuracy: 0.8999 - val_loss: 0.3398 - val_accuracy: 0.8770
Epoch 8/10
60000/60000 [=====] - 5s 80us/step - loss: 0.2557
- accuracy: 0.9046 - val_loss: 0.3461 - val_accuracy: 0.8737
Epoch 9/10
60000/60000 [=====] - 5s 85us/step - loss: 0.2450
- accuracy: 0.9086 - val_loss: 0.3459 - val_accuracy: 0.8791
Epoch 10/10
60000/60000 [=====] - 5s 78us/step - loss: 0.2366
- accuracy: 0.9113 - val_loss: 0.3466 - val_accuracy: 0.8854
```

```
Out[6]: <keras.callbacks.callbacks.History at 0x17b3c105188>
```

Step 4: Evaluating the model with Testing dataset

```
In [7]: 1 scores = model_T1.evaluate(x_test, y_test, verbose=1)
        2 print("%s: %.2f%%" % (model_T1.metrics_names[1], scores[1]*100))
```

```
10000/10000 [=====] - 0s 33us/step
accuracy: 88.54%
```

Step 5: Changing the Model Structure to get better evaluation results

Changing the number of neurons:

```
In [8]: 1 model_T1 = Sequential()
2 model_T1.add(Flatten(input_shape=(28, 28, 1)))
3 model_T1.add(Dense(256, activation='relu'))
4 model_T1.add(Dense(128, activation='relu'))
5 model_T1.add(Dense(10, activation='softmax'))
6
7 # Defining Optimizer, Loss function and evaluation metrics
8 model_T1.compile(loss='categorical_crossentropy', optimizer='adam', met
9
10 # Setting a random state so that comparing the evaluation metrics is po
11 tf.set_random_seed(1)
12 np.random.seed(1)
13 model_T1.fit(x_train, y_train, epochs=10, batch_size=32, validation_dat
14
15 scores = model_T1.evaluate(x_test, y_test, verbose=1)
16 print("%s: %.2f%%" % (model_T1.metrics_names[1], scores[1]*100))
```

```
10000/10000 [=====] - 1s 51us/step
accuracy: 88.70%
```

Task 2: Build a neural network with the use of convolutional layers

Step 1: Defining Model Architecture

```
In [9]: 1 model_T2 = Sequential()
2 #Convolutional Layer 1
3 model_T2.add(Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28,
4 padding='same'))
5 #Pooling Layer 1
6 model_T2.add(MaxPooling2D((2, 2)))
7 #Convolutional Layer 2
8 model_T2.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
9 #Pooling Layer 1
10 model_T2.add(MaxPooling2D((2, 2)))
11 #Flatten Layer
12 model_T2.add(Flatten())
13 #Fully-Connected Layer 1
14 model_T2.add(Dense(128, activation='relu'))
15 #Fully-Connected Layer 2 (Output Layer)
16 model_T2.add(Dense(10, activation='softmax'))
```

```
WARNING:tensorflow:From C:\Users\nikhi\anaconda3\envs\NeuralNetwork\lib\si
te-packages\keras\backend\tensorflow_backend.py:4070: The name tf.nn.max_p
ool is deprecated. Please use tf.nn.max_pool2d instead.
```

Step 2: Compiling the Model

```
In [10]: 1 # Defining Optimizer, Loss function and evaluation metrics
2 from keras.optimizers import Adam
3 custom_adam = Adam(lr=0.002)
4
5 model_T2.compile(loss='categorical_crossentropy', optimizer=custom_adam)
```

```
In [11]: 1 # Model Structure
2 model_T2.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d_1 (MaxPooling2)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 14, 14, 64)	18496
max_pooling2d_2 (MaxPooling2)	(None, 7, 7, 64)	0
flatten_3 (Flatten)	(None, 3136)	0
dense_7 (Dense)	(None, 128)	401536
dense_8 (Dense)	(None, 10)	1290
=====		
Total params: 421,642		
Trainable params: 421,642		
Non-trainable params: 0		
=====		

Step 3: Training the Model

```
In [12]: 1 # Setting a random state so that comparing the evaluation metrics is po
2         tf.set_random_seed(1)
3         np.random.seed(1)
4
5         model_T2.fit(x_train, y_train, epochs=10, batch_size=32, validation_dat
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 58s 959us/step - loss: 0.38
01 - accuracy: 0.8636 - val_loss: 0.3107 - val_accuracy: 0.8880

Epoch 2/10

60000/60000 [=====] - 58s 969us/step - loss: 0.24
97 - accuracy: 0.9080 - val_loss: 0.2640 - val_accuracy: 0.9023

Epoch 3/10

60000/60000 [=====] - 58s 972us/step - loss: 0.20
82 - accuracy: 0.9222 - val_loss: 0.2709 - val_accuracy: 0.9051

Epoch 4/10

60000/60000 [=====] - 59s 991us/step - loss: 0.17
69 - accuracy: 0.9334 - val_loss: 0.2417 - val_accuracy: 0.9097

Epoch 5/10

60000/60000 [=====] - 58s 966us/step - loss: 0.14
91 - accuracy: 0.9435 - val_loss: 0.2567 - val_accuracy: 0.9156

Epoch 6/10

60000/60000 [=====] - 59s 979us/step - loss: 0.12
69 - accuracy: 0.9513 - val_loss: 0.2710 - val_accuracy: 0.9087

Epoch 7/10

60000/60000 [=====] - 58s 967us/step - loss: 0.10
72 - accuracy: 0.9589 - val_loss: 0.2906 - val_accuracy: 0.9142

Epoch 8/10

60000/60000 [=====] - 59s 985us/step - loss: 0.09
22 - accuracy: 0.9656 - val_loss: 0.3263 - val_accuracy: 0.9092

Epoch 9/10

60000/60000 [=====] - 59s 983us/step - loss: 0.07
64 - accuracy: 0.9709 - val_loss: 0.3546 - val_accuracy: 0.9145

Epoch 10/10

60000/60000 [=====] - 59s 983us/step - loss: 0.07
08 - accuracy: 0.9730 - val_loss: 0.4068 - val_accuracy: 0.9114

Out[12]: <keras.callbacks.callbacks.History at 0x17b3c5a31c8>

Step 4: Evaluating the model with Testing dataset

```
In [13]: 1 scores = model_T2.evaluate(x_test, y_test, verbose=1)
2         print("%s: %.2f%%" % (model_T1.metrics_names[1], scores[1]*100))
```

10000/10000 [=====] - 2s 232us/step
accuracy: 91.14%

Step 5: Changing the Model Structure to get better evaluation results

Changing the model structure by increasing the number of filters in the convolutional layer to 64 and 128 respectively

```
In [14]: 1 model_T2 = Sequential()
2 model_T2.add(Conv2D(64, (3, 3), activation='relu', input_shape=(28, 28,
3 padding='same')))
4 model_T2.add(MaxPooling2D((2, 2)))
5 model_T2.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
6 model_T2.add(MaxPooling2D((2, 2)))
7 model_T2.add(Flatten())
8 model_T2.add(Dense(128, activation='relu'))
9 model_T2.add(Dense(10, activation='softmax'))
10
11 from keras.optimizers import Adam
12 custom_adam = Adam(lr=0.002)
13
14 model_T2.compile(loss='categorical_crossentropy', optimizer=custom_adam)
15
16 # Setting a random state so that comparing the evaluation metrics is po
17 tf.set_random_seed(1)
18 np.random.seed(1)
19 model_T2.fit(x_train, y_train, epochs=10, batch_size=32, validation_dat
20
21 scores = model_T2.evaluate(x_test, y_test, verbose=1)
22 print("%s: %.2f%%" % (model_T1.metrics_names[1], scores[1]*100))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 125s 2ms/step - loss: 0.367

1 - accuracy: 0.8677 - val_loss: 0.2994 - val_accuracy: 0.8927

Epoch 2/10

60000/60000 [=====] - 122s 2ms/step - loss: 0.240

9 - accuracy: 0.9103 - val_loss: 0.2551 - val_accuracy: 0.9085

Epoch 3/10

60000/60000 [=====] - 126s 2ms/step - loss: 0.197

6 - accuracy: 0.9259 - val_loss: 0.2677 - val_accuracy: 0.9026

Epoch 4/10

60000/60000 [=====] - 128s 2ms/step - loss: 0.164

6 - accuracy: 0.9387 - val_loss: 0.2484 - val_accuracy: 0.9119

Epoch 5/10

60000/60000 [=====] - 123s 2ms/step - loss: 0.137

7 - accuracy: 0.9483 - val_loss: 0.2487 - val_accuracy: 0.9123

Epoch 6/10

60000/60000 [=====] - 127s 2ms/step - loss: 0.114

8 - accuracy: 0.9554 - val_loss: 0.2977 - val_accuracy: 0.9116

Epoch 7/10

60000/60000 [=====] - 127s 2ms/step - loss: 0.093

4 - accuracy: 0.9647 - val_loss: 0.3319 - val_accuracy: 0.9152

Epoch 8/10

60000/60000 [=====] - 125s 2ms/step - loss: 0.079

5 - accuracy: 0.9700 - val_loss: 0.3438 - val_accuracy: 0.9174

Epoch 9/10

60000/60000 [=====] - 123s 2ms/step - loss: 0.068

4 - accuracy: 0.9753 - val_loss: 0.3628 - val_accuracy: 0.9175

Epoch 10/10

60000/60000 [=====] - 123s 2ms/step - loss: 0.061

6 - accuracy: 0.9770 - val_loss: 0.4392 - val_accuracy: 0.9158

10000/10000 [=====] - 5s 455us/step

accuracy: 91.58%

Task 3: Change the type of optimizer or learning rate that you applied in the previous tasks, and see how these changes can influence model performance

In [15]:

```
1 #Testing custom sgd with different learning rates
2 lrate = (0.002, 0.004, 0.006, 0.008)
3 for lr in lrate:
4     epochs = 10
5     decay = lr/epochs
6     sgd = SGD(lr=lr, momentum=0.7, decay=decay, nesterov=False) #Stochastic Gradient Descent
7     # Compile model
8     model_T2.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
9     model_T2.fit(x_train, y_train, epochs=epochs, batch_size=32, validation_data=(x_test, y_test))
10    scores = model_T2.evaluate(x_test, y_test, verbose=1)
11    print("learning rate: ", lr)
12    print("%s: %.2f%%" % (model_T2.metrics_names[1], scores[1]*100))
```

10000/10000 [=====] - 5s 464us/step

learning rate: 0.002

accuracy: 92.65%

10000/10000 [=====] - 6s 600us/step

learning rate: 0.004

accuracy: 92.74%

10000/10000 [=====] - 4s 426us/step

learning rate: 0.006

accuracy: 92.73%

10000/10000 [=====] - 8s 834us/step

learning rate: 0.008

accuracy: 92.80%